

## I. Agent Structure (MADDPG - Multi-Agent Deep Deterministic Policy Gradient)

**agents** - list of all the DDPG agents (the number of agents is equal to the number of players in the game)

Details of DDPG is in the [project 2 report](#)

### Parameters

num_agents	2	Number agents (players)
n_episodes	1000	Number of times the environment is run
BUFFER_SIZE	1e6	Replay (memory) buffer size
BATCH_SIZE	128	Minibatch size
GAMMA	0.99	Discount factor
TAU	1e-3	Soft update of target parameters
LR_ACTOR	1e-4	Learning rate of the actor
LR_CRITIC	1e-4	Learning rate of the critic
WEIGHT_DECAY	0	L2 weight decay
mu	0	Noise parameter
theta	0.05	Noise parameter
sigma	0.4	Noise parameter

### Overview

For each episode:

- 1) Reset the environment
- 2) Get the state of the environment

While the environment is running:

- 1) Use the agent to **predict action** based on the environment's state
- 2) Pass the predict action into the environment
- 3) Retrieve the next state, reward, and done variable
- 4) **Agent step**. Record next state, reward, and done variable to **memory** in agent
- 5) Set the current state variable to next state
- 6) Add reward to score
- 7) If the current score is the highest, save the model
- 8) Break out of while loop if done is true

### Predict action

- 1) Loop through the states that were passed in


- 2) For each state, pass it to the corresponding agent
- 3) Collect the action and return the actions as an array

#### **Agent step**

- 1) Loop through the states that were passed in
- 2) For each state, actions, rewards, next\_states, dones, pass it into the corresponding agent

## **II. Results**

Here's a graph of the episodes vs scores. Most scores were between 1 and 1.5 except for a 2.5. The yellow line is the average score. The highest average score is around 0.75



## **III. Ideas for Future work**

Other improvements to the algorithm I would like to pursue:

- 1) Prioritized Experience Replay
- 2) Experiment with sharing the replay buffer with all agents